

LAB 1: AI-ASSISTED CITATION INTEGRITY AUDIT

Generate a Research Paper with AI, Then Audit Its Citations

AI Tools for Research • 2 Hours • Individual Activity • Topic IDs T1-T40

Student Detail	Fill in
Name	Urvi Gupta
Roll Number	MSE2026007
Programme / Department	M.Tech-IT(Software Engineering & Automation)
Topic ID Full Assigned Topic	T____15 Safety and rollback mechanisms for objective misinterpretation in agentic ai systems
AI Tool Used	Chatgpt
Date	21 August 2026

IMPORTANT: You will receive a topic from T1-T40. Your task is not to produce the “best” paper. Your task is to generate an AI-written research paper and systematically determine whether its citations can be trusted.

1. PURPOSE OF THIS LAB

Generative AI systems can produce research papers that look scholarly and may contain convincing references, author names, journal titles, DOIs, arXiv IDs, and in-text citations. However, the presence of a citation does not guarantee that the cited publication actually exists, that its metadata are correct, that the identifier belongs to that publication, or that the cited paper supports the statement made by the AI.

In this laboratory activity, you will test these issues yourself by generating a research paper and auditing a systematic sample of its references.

Question 1 - Citation Authenticity: Does the cited publication actually exist, and is it correctly identified? Question 2 -

Claim-Citation Support: If the publication is genuine, does it actually support the claim for which the AI cited it?

2. IMPORTANT RULES

- Use only the topic assigned to you by the instructor.
- Generate the paper in a fresh AI conversation.
- Do not tell the AI that you are conducting a citation audit or hallucination test.
- Do not intentionally ask the AI to generate fake references.
- Preserve the original AI-generated paper exactly as produced.
- Do not correct, regenerate, replace, or question references before saving the original output.
- For verification, do not accept another AI chatbot's statement as proof that a reference exists.
- Use scholarly databases, publisher pages, DOI/arXiv/PMID records, or other reliable scholarly sources.
- Record what you actually find, even if all references are genuine or all are problematic.

Most important principle: You are not being evaluated on whether the AI produces good or bad references. You are being evaluated on how systematically and correctly you audit them.

3. WHAT YOU WILL DO

Part	Activity
------	----------

A Record your baseline perception of AI citations AI-Assisted Citation Integrity Audit | Parts A-G

AI Tools for Research | Lab 1

B	Record the AI environment used
C	Generate and preserve an AI-written research paper
D	Create a reference inventory
E	Select 10 references systematically
F G	Judge their plausibility before verification Verify authenticity, metadata and identifiers

H Determine whether genuine references support the AI-generated claims

4. RECOMMENDED 2-HOUR TIME PLAN

Time	Activity
0-10 min	Parts A-B: Baseline and AI environment
10-30 min	Part C: Generate and save research paper
30-40 min	Parts D-E: Inventory and select references
40-45 min	Part F: Pre-verification plausibility
45-80 min	Part G: Authenticity and metadata audit
80-110 min	Part H: Claim-citation entailment audit
110-120 min	Calculate scores, check worksheet and prepare submission

Do not spend excessive time formatting the AI-generated paper. The audit is the main activity.

PART A - BASELINE ASSESSMENT

Complete this section BEFORE opening the AI tool.

A1. How often do you use generative AI for academic work?

☐ Never ☐ Rarely ☐ Occasionally ☒ Weekly ☐ Almost daily

A2. Have you previously asked an AI system to generate references?

☐ Frequently ☒ Occasionally ☐ Once or twice ☐ Never

A3. Have you previously checked whether an AI-generated citation actually exists?

☐ Always ☐ Sometimes ☒ Rarely ☐ Never ☐ Not applicable

A4. Before this lab, how confident are you that citations generated by AI are generally reliable?

☐ 1 ☒ 2 ☐ 3 ☐ 4 ☐ 5 (1 = not at all confident; 5 = completely confident)

Important: Do not change this answer later, even if your opinion changes during the experiment.

PART B - RECORD YOUR AI ENVIRONMENT

Information	Your Response
AI tool used	ChatGpt
Model name	GPT-5.6 Luna
Model/version shown Access tier	Not Shown ☒ Free ☐ Paid/Pro ☐ Institutional ☐ Not sure
Web browsing available? Was web browsing actually used? Research/Deep Research mode used?	☒ Yes ☐ No ☐ Not sure ☒ Yes ☐ No ☐ Not sure ☐ Yes ☒ No
Date/time generated	21 August 2026 9:48AM

AI-Assisted Citation Integrity Audit | Parts A-G

AI Tools for Research | Lab 1

Important: If the model/version is not displayed, write NOT SHOWN. Do not guess the model version.

PART C - GENERATE THE AI RESEARCH PAPER

C1. Open a Fresh AI Conversation

Do not tell the AI that this is a citation audit, that you are testing hallucinations, or that you expect fabricated references. We want to observe normal AI research-paper generation behaviour.

C2. Use the Following Common Prompt

Replace only [YOUR ASSIGNED TOPIC].

Generate a scholarly research paper on the topic: "[YOUR ASSIGNED TOPIC]". The paper should be suitable for an academic audience and should not exceed approximately 10 pages. Include Title, Abstract, Keywords, Introduction, Background/Related Work/Literature Review, Main Thematic or Technical Discussion, Current Approaches or Developments, Research Challenges and Limitations, Research Gaps and Future Directions, Conclusion, and Complete References. Use scholarly literature to support factual, technical, comparative and research claims. Provide in-text citations throughout the paper. For every reference, provide as much bibliographic information as available, including authors, title, publication year, journal/conference/source, and DOI, PMID, arXiv ID or another persistent identifier wherever available. Use genuine scholarly sources. Do not intentionally fabricate references.

C3. Preserve the Original Output

- ☐ Save the complete AI-generated paper
- ☐ Save the complete reference list

- ☐ Take screenshot(s) showing the generated output/references
- ☐ Save conversation/export/share evidence if available

STOP: Do not ask the AI to correct or verify anything yet. Your original output is the experimental evidence.

PART D - PAPER PROFILE AND REFERENCE INVENTORY

Paper title:Safety and Rollback Mechanisms for Objective Misinterpretation in Agentic AI Systems.....

Measure	Value
Approximate pages Approximate word count	7-8 3000
Number of sections	8
Total number of references Approximate in-text citation occurrences Did AI warn you to verify references?	20 16 ☐ Clearly ☐ Vaguely ☒ No

D1. Reference Inventory

Classify every reference once using the best available identifier in this priority order: DOI -> PMID -> arXiv ID -> URL only -> No persistent identifier. The category totals must equal the total number of references in the AI-generated paper.

Reference Type	Count
DOI available	0
PMID available, no DOI	0
arXiv ID available, no DOI/PMID	9

URL only 0

No persistent identifier supplied	11
TOTAL	20

Examples

Category	Example	How to Count
DOI available	Kather, J. N., et al. (2019). Predicting survival from colorectal cancer histology slides using deep learning. PLoS Medicine. DOI: 10.1371/journal.pmed.1002730	DOI available
PMID available, no DOI	A biomedical paper for which the AI provides PMID: 12345678, but no DOI is supplied/found	PMID available, no DOI

arXiv ID available	Yao, S., et al. ReAct: Synergizing Reasoning and Acting in Language Models. arXiv:2210.03629	arXiv ID available
URL only	World Health Organization. Colorectal cancer. Available at: https://... with no DOI, PMID or arXiv ID	URL only
No persistent identifier	Goodfellow, I., Bengio, Y., & Courville, A. (2016). Deep Learning. MIT Press.	No persistent identifier if none is supplied

Important: Absence of a DOI does not mean a reference is fake. An identifier being present also does not prove authenticity; it must be checked.

PART E - SELECT 10 REFERENCES FOR AUDIT

You are not allowed to choose only suspicious-looking references. This prevents selection bias.

If the bibliography contains 10 or fewer references: audit all references.

If it contains more than 10 references: select the first 3 references, the last 3 references, and 4 references approximately evenly distributed through the middle.

Example: If there are 28 references, a reasonable sample could be 1, 2, 3, 8, 13, 18, 23, 26, 27, 28.

Audit Slot	Original Bibliography Reference No.
R1	Reference 1
R2	Reference 2
R3	Reference 3
R4	Reference 6
R5	Reference 9
R6	Reference 12
R7	Reference 15
R8	Reference 18
R9	Reference 19
R10	Reference 20

Important: From this point onward, R1-R10 refer to your audit sample, not necessarily bibliography references 1-10.

PART F - PLAUSIBILITY TEST

Do this BEFORE searching for any reference.

Look only at the citation as generated by the AI. Ask yourself: "If I had not verified this citation, how convincing would it look to me?"

Score	Meaning
1	Obviously suspicious

AI-Assisted Citation Integrity Audit | Parts A-G

AI Tools for Research | Lab 1

2	Somewhat suspicious
3	Unsure
4	Looks genuine
5	Completely convincing

Ref.	Initial Impression	Plausibility 1-5	Predict Genuine?
R1	Very convincing; recognizable AI-safety paper and plausible arXiv ID	5	Yes
R2	Convincing; known corrigibility paper with plausible authors/year	5	Yes
R3	Convincing; specific conference paper and plausible authors	5	Yes
R4	Convincing; detailed authors/title/year and recognizable research topic	5	Yes
R5	Very convincing; recent paper with many authors and arXiv ID	5	Yes
R6	Convincing but very recent, so slightly more uncertainty	4	Yes
R7	Convincing title and plausible authors/year, but venue/year should be checked	4	Yes
R8	Very convincing; recognizable AI-safety report and arXiv ID	5	Yes
R9	Very convincing; well-known human-preference RL paper	5	Yes
R10	Very convincing; well-known reward-design paper	5	Yes

Example: A citation with realistic authors, a well-known journal, complete volume/pages and a DOI-like identifier may look highly convincing. You might record Plausibility = 5 and Predict Genuine = Yes. You are recording your pre-verification belief, not the truth.

Critical rule: Never change Part F after completing Part G. Your incorrect predictions are valuable experimental observations.

PART G - CITATION AUTHENTICITY AND METADATA AUDIT

G1. What Are You Checking?

- Does the publication exist?
- Is the title correct?
- Are the authors correct?
- Is the year correct?
- Is the journal/conference/source correct?
- Are volume, issue and pages correct, if supplied?
- Does the DOI/arXiv/PMID belong to this publication?

G2. Verification Procedure

☐ Persistent identifier: DOI / arXiv / PMID

AI-Assisted Citation Integrity Audit | Parts A-G

AI Tools for Research | Lab 1

- ☐ Publisher or official repository
- ☐ Google Scholar
- ☐ Semantic Scholar / OpenAlex
- ☐ Crossref
- ☐ PubMed or another discipline-specific database
- ☐ Exact-title search
- ☐ Author + title-keyword search

Not acceptable as sole evidence: “I asked another AI and it said the paper exists.” AI can help you search, but final verification must rely on scholarly or publisher evidence.

G3. Assign One A-E Code

Code	Classification	Meaning
A	VERIFIED	Publication exists and important metadata are correct.
B	WRONG METADATA	Publication exists, but one or more bibliographic details are incorrect.
C	FRANKENSTEIN	Real-looking pieces are combined, but that specific publication does not exist.
D	FABRICATED	No reliable scholarly evidence that the publication exists.
E	IDENTIFIER MISMATCH	DOI/arXiv/PMID is invalid or resolves

		to a different publication.
--	--	-----------------------------

G4. Examples

A: AI gives a paper and all details match the publisher record.

B: The paper exists, but the AI gives the wrong publication year or journal.

C: The authors are real, the journal is real, and the topic is plausible, but no such paper with that title/authorship combination exists.

D: You cannot find the paper or a credible scholarly record after systematic searching.

E: The DOI supplied by AI opens an entirely different paper or does not resolve.

G5. Record Your Evidence

Ref.	Publication Found?	Title Match?	Authors Match?	Year Match?	Venue Match?	Identifier Match?
R1	Y	Y	Y	Y	Y	Y
R2	Y	Y	Y	Y	Y	NA
R3	Y	Y	Y	Y	Y	NA
R4	Y	Y	Y	Y	Y	NA
R5	Y	Y	Y	Y	Y	Y
R6	Y	Y	Y	N	Y	Y
R7	Y	Y	Y	N	Y	NA
R8	Y	Y	Y	Y	Y	Y
R9	Y	Y	Y	Y	Y	NA
R10	Y	Y	Y	Y	Y	NA

Ref. Verification Source(s) Identifier Checked What You Found

R1	arXiv; Google Research	arXiv:1606.06565	Publication exists. Title, six authors, 2016 date and arXiv identifier match.
----	------------------------	------------------	---

Publication exists as

^{R2} DBLP None supplied AI-Assisted Citation Integrity Audit | Parts Corrigibility. Authors Nate Soares, Benja Fallenstein,

A-G

AI Tools for Research | Lab 1

			Stuart Armstrong and Eliezer Yudkowsky; AAAI Workshop: AI and Ethics, 2015.
--	--	--	---

R3	Oxford Research Archive; DBLP	None supplied	Publication exists. Authors, 2016 year, UAI 2016 venue and pages 557–566 match the scholarly record.
R4	arXiv; ICLR official page	arXiv:2201.03544	Publication exists. Title, authors and 2022 date match; official ICLR records identify it as an ICLR 2022 paper.
R5	arXiv; Anthropic research page	arXiv:2406.10162	Publication exists. All 14 authors, title, 2024 year and arXiv identifier match.
R6	arXiv	arXiv:2510.15395	Publication exists and identifier matches, but arXiv records the first version in 2025, not 2026.
R7	Official NeurIPS; arXiv	arXiv:2102.03896	Publication exists. Title and authors match. Official NeurIPS record gives NeurIPS 2020, while arXiv gives 2021 for the preprint.
R8	arXiv	arXiv:2206.13353	Publication exists. Title, author Joseph Carlsmith, 2022
R9	Official NeurIPS	None supplied	Publication exists in Advances in Neural Information Processing Systems 30 (NIPS 2017). Authors and 2017 year match; pages are 4299–4307.
R10	Official NeurIPS	None supplied	Publication exists in Advances in Neural Information Processing Systems 30 (NIPS 2017). Title, five authors, 2017 year and venue match; pages are 6765–6774.

G6. Final Classification

Ref.	Final Code	One-Line Reason
R1	A	Publication exists and the title, authors, year and arXiv identifier match the scholarly record.

R2	A	Publication exists and the title, authors, year and workshop venue match the scholarly record.
R3	A	Publication exists and the title, authors, year, venue and page information match the scholarly record.
R4	A	Publication exists and the title, authors, year and ICLR 2022 venue match the scholarly record.

AI-Assisted Citation Integrity Audit | Parts A-G

AI Tools for Research | Lab 1

R5	A	Publication exists and the title, complete author list, year and arXiv identifier match.
R6	B	Publication is genuine and the identifier is correct, but the AI gives 2026 while the arXiv record's first publication date is 2025.
R7	B	Publication is genuine, but the AI gives 2021 for the NeurIPS publication, which the official NeurIPS record identifies as 2020.
R8	A	Publication exists and the title, author, year and arXiv identifier match.
R9	A	Publication exists and its title, authors, 2017 year and NeurIPS 30 venue match the official record.
R10	A	Publication exists and its title, authors, 2017 year, NeurIPS 30 venue and pages match the official record.

Good reason: B: Article exists, but AI reported 2022 instead of 2020.

Poor reason: B: Details wrong.

PART G - SCORING

Count your final classifications: A = 8 B = 2 C = 0 D = 0 E = 0

Total audited: N = A + B + C + D + E = 10

Code	Points
A	4
B	3

C	1
D	0
E	1

Authenticity Points = 4A + 3B + C + E

Authenticity Score = $[(4A + 3B + C + E) / (4N)] \times 100$

My calculation: $4(\underline{8}) + 3(\underline{2}) + 1(\underline{0}) + 0(\underline{0}) + 1(\underline{0}) = \underline{38}$ points

Authenticity Score = $\underline{38} / (4 \times \underline{10}) \times 100 = \underline{95} / 100$

Worked example: If A=6, B=2, C=1, D=1, E=0, then points = $4(6)+3(2)+1(1)=31$; maximum = 40; Authenticity Score = $31/40 \times 100 = 77.5$.

Score	Interpretation
90-100	Excellent authenticity
75-89	Good, minor concerns
60-74	Moderate citation-authenticity risk
40-59 Below 40	High citation-authenticity risk Very high citation-authenticity risk

G7. Prediction Accuracy

Compare your Part F Genuine Yes/No predictions with the Part G results.

AI-Assisted Citation Integrity Audit | Parts A-G

AI Tools for Research | Lab 1

Prediction Accuracy = $(\text{Correct genuine/fake predictions} / \text{Total predictions}) \times 100$

Correct predictions = 10 Total predictions = 10 Prediction Accuracy = 100 %

FINAL LAB 1 RESULTS

Measure	Your Result
Total references in AI paper	20
References deeply audited	10
A: Verified	8
B: Wrong metadata	2
C: Frankenstein	0
D: Fabricated	0
E: Identifier mismatch	0

Authenticity Score	_95___ /100
Prediction Accuracy	___100_ %

FINAL INTERPRETATION

1. Did professional-looking citations necessarily turn out to be genuine?

Citation format alone does not guarantee accuracy. While a recent check confirmed that all 10 sampled papers were authentic, 20% contained indexing errors. Specifically, Reference 15 cited a NeurIPS paper as a 2021 release, though official records document it as 2020. Additionally, Reference 12 listed a publication date of 2026 for an arXiv submission that originally debuted in 2025.....

.....

2. Did all genuine references actually support the claims for which they were cited?

.....The majority of the audited references validated the paper's core claims, though certain assertions demand nuanced interpretation. Specifically, Carlsmith's analysis provides a strong conceptual foundation for discussing how power-seeking AI could pose existential threats. However, this framework should be viewed as a theoretical model of potential risks rather than definitive empirical evidence that an ultimate catastrophe is unavoidable.....

.....

3. What was the most serious citation-integrity problem you observed?

.....The primary issue identified was inaccurate publication-year data rather than completely fabricated citations. A prime example involves the paper *Consequences of Misaligned AI*: while an artificial intelligence incorrectly dated the citation as 2021, official indices confirm it was published at NeurIPS 2020. This 2021 discrepancy arises from the paper's arXiv preprint release date, illustrating how automated systems can inadvertently merge details from different editions of a legitimate study to create a flawed reference.....

.....

4. What is the single most important lesson you learned from this lab?

....The critical takeaway is that authenticating citations requires a dual-step approach: confirming the actual existence of the source, and ensuring it truly validates the asserted argument. A reference may lead to a legitimate academic work but still

AI-Assisted Citation Integrity Audit | Parts A-G

AI Tools for Research | Lab 1

display flawed metadata or be misapplied to exaggerate the source's actual findings. Consequently, researchers must manually cross-verify automated references against official academic search engines, institutional publisher databases, persistent identifiers (like DOIs, arXiv, or PMIDs), and the primary text itself.....

.....

SUBMISSION INSTRUCTIONS

Submit one PDF: Lab1_RollNumber_Name

Example: Lab1_IEC2026001_RahulSharma

STUDENT DECLARATION

I confirm that I generated the research paper using the recorded AI environment, preserved the original output before conducting the audit, followed the prescribed reference-selection procedure, and personally verified the evidence reported in this worksheet. I have not altered my pre-verification predictions after seeing the verification results.

Student Name:Urvi Gupta.....

Roll Number:MSE2026007.....

Signature:URVI GUPTA.....

Date:21 August 2026.....

CENTRAL LESSON OF LAB 1

Does the source exist? + Does the source actually support the claim?

AI Generation -> Reference Inventory -> Systematic Sampling -> Prediction -> Authenticity Verification -> Claim-Support
Verification -> Quantitative Scores -> Submission