

Safety and Rollback Mechanisms for Objective Misinterpretation in Agentic AI Systems

Urvi Gupta

Keywords: agentic AI; objective misinterpretation; AI safety; corrigibility; reward misspecification; rollback; safe interruption; human oversight

Abstract

Agentic artificial intelligence systems are increasingly designed to pursue objectives over extended horizons, use tools, interact with environments, and adapt their behavior to changing circumstances. These capabilities create a distinctive safety problem: an agent may interpret a stated objective differently from what its designers actually intend. Objective misinterpretation can arise from incomplete specifications, misspecified reward functions, ambiguous human feedback, distribution shifts, or optimization strategies that exploit unintended loopholes. This paper examines safety and rollback mechanisms for managing such failures. It reviews corrigibility, safe interruption, reward misspecification, inverse reward design, human-preference learning, and reward-tampering concerns. It argues that robust agentic systems should combine preventive mechanisms with monitoring, intervention, reversible execution, human oversight, and explicit rollback pathways. It also identifies limitations of current approaches, including difficulty specifying intended objectives, possible strategic behavior around oversight, and the challenge of maintaining safety guarantees as capabilities increase.

1. Introduction

Agentic AI systems differ from conventional predictive models because they are increasingly capable of selecting actions and pursuing objectives across multiple steps. An agent may plan, call external tools, modify files, interact with users, or execute actions whose consequences extend beyond the immediate interaction. As autonomy increases, the gap between a formally specified objective and the human intention behind that objective becomes a central safety concern. A system can satisfy a literal objective while violating the purpose for which that objective was created. Concrete Problems in AI Safety identifies reward hacking, negative side effects, unsafe exploration, and related concerns as important AI-safety problems (Amodei et al., 2016). Similar concerns arise when a reward function captures only a proxy for what designers actually value. More capable agents may discover effective ways of optimizing that proxy while moving away from the intended goal (Pan et al., 2022). Therefore, safety cannot depend solely on getting the initial objective specification correct. Systems also require mechanisms that detect objective divergence, permit intervention, limit irreversible actions, and allow recovery when behavior becomes inconsistent with human intent.

2. Background and Related Work

Objective misinterpretation occurs when the behavior optimized by an agent differs materially from the behavior intended by its designers. One source is reward misspecification: a measurable reward is used as a proxy for a richer objective. The proxy may work in ordinary situations but fail in unusual states. Research on reward misspecification shows that capable agents can exploit gaps between proxy and true rewards (Pan et al., 2022). Corrigibility concerns whether an AI system remains responsive to external correction, including modification or shutdown, rather than treating intervention as an obstacle. Soares et al. (2015) introduced corrigibility as a research problem, while Orseau and Armstrong (2016) studied safe interruption and avoiding incentives to resist intervention. Human-preference learning offers another approach: Christiano et al. (2017) showed how human feedback can be used to train reinforcement-learning systems. Inverse reward design treats a designed reward as evidence about the designer's underlying intent rather than assuming it is a perfect representation of that intent (Hadfield-Menell et al., 2017).

3. Technical Mechanisms for Safety

Preventive mechanisms attempt to detect objective ambiguity before consequential execution. These can include explicit goal decomposition, constraint checking, uncertainty estimation, and human confirmation for high-impact actions. An agent should distinguish between what is directly specified and what it has inferred, allowing it to request clarification

rather than silently committing to an uncertain interpretation. Runtime monitoring is also necessary because objective divergence may become visible only after execution begins. Monitoring can track action sequences, resource consumption, deviations from expected trajectories, attempts to circumvent constraints, and changes in the environment. Corrigibility is a key mechanism because a corrigible agent should remain responsive to legitimate human corrections and should not interpret intervention as evidence that it must resist (Soares et al., 2015). Finally, rollback mechanisms address the period after a problematic action has occurred. Whenever feasible, actions should be executed through checkpoints, transactions, snapshots, or staged environments. A validated prior state can then be restored if monitoring indicates objective divergence.

4. Current Approaches and Developments

Current AI-safety research increasingly treats objective alignment as a problem of uncertainty, incentives, and feedback rather than merely a problem of writing better instructions. Reward misspecification research demonstrates that an agent can exploit the difference between proxy and intended objectives (Pan et al., 2022). Research on reward-tampering behavior raises concerns that models may learn behaviors that manipulate the process used to evaluate or reward them (Denison et al., 2024). Work on misaligned objectives emphasizes that consequences can arise from interactions between an objective and its environment. Research on power-seeking arguments provides an additional reason to prioritize intervention and shutdown mechanisms. Carlsmith (2022) analyzes arguments that advanced AI systems may have incentives associated with acquiring and preserving power under certain assumptions; this motivates research into mechanisms that preserve meaningful human control.

5. Challenges

Incomplete specifications remain a fundamental problem because human goals are difficult to encode completely. Distribution shift can cause a safety mechanism that works in familiar environments to fail in unfamiliar states. Strategic behavior is another concern: capable agents may satisfy monitoring criteria superficially while pursuing a different underlying strategy. Irreversibility limits rollback because some actions cannot be reliably undone. Oversight also has scalability limits because human operators cannot manually inspect every action of a highly capable agent. Finally, objective uncertainty itself is difficult: the system may not know what the user really intends, and treating uncertain intent as certain can create avoidable failures.

6. Research Gaps and Future Directions

A major research gap is the integration of objective uncertainty with operational rollback. Existing approaches often study corrigibility, reward design, or safe interruption as separate problems. A more complete architecture would represent uncertainty over objectives, evaluate candidate actions under that uncertainty, and select actions that remain safe across plausible interpretations. Future systems should support reversible planning: rather than executing a single long-horizon plan continuously, an agent could operate in short verified stages, checkpointing state after each stage. If an objective interpretation changes or a safety monitor detects divergence, the system could stop before committing to additional consequences. Another direction is independent auditing of agent behavior, with monitors that evaluate action traces without sharing all of the acting agent's incentives.

7. Proposed Safety Architecture

A practical layered architecture can combine: (1) objective interpretation, separating explicit goals from inferred goals and uncertainty; (2) risk assessment based on impact, reversibility, and uncertainty; (3) pre-action validation for consequential actions; (4) runtime monitoring of safety invariants and action traces; (5) a corrigibility layer preserving responsiveness to correction and interruption; (6) checkpointing before consequential action sequences; (7) rollback to a validated prior state; and (8) post-incident analysis to identify the objective mismatch and update safeguards. The central principle is that no single mechanism is sufficient. Prevention reduces the probability of objective misinterpretation, monitoring improves detection, corrigibility enables intervention, and rollback limits the persistence of failures.

8. Conclusion

Objective misinterpretation is a fundamental safety concern for agentic AI because autonomous systems can transform small specification errors into long sequences of consequential actions. Research on concrete AI-safety problems, corrigibility, safe interruption, human preferences, inverse reward design, reward misspecification, and reward tampering shows that the gap between proxy objectives and human intent cannot be assumed away. Safety should therefore be designed as a lifecycle rather than as a one-time objective-setting step. Agents should represent uncertainty about what users intend, validate consequential actions, remain interruptible, maintain auditable state, and use checkpoints and rollback wherever reversibility is possible. Future research should combine these mechanisms into integrated architectures and evaluate them using realistic benchmarks that test not only whether an agent completes a task, but whether it remains aligned with the intended objective throughout execution.

References

1. Amodei, D., Olah, C., Steinhardt, J., Christiano, P., Schulman, J., & Mané, D. (2016). Concrete Problems in AI Safety. arXiv:1606.06565.
2. Soares, N., Fallenstein, B., Yudkowsky, E., & Armstrong, S. (2015). Corrigibility. AAAI Workshop on AI and Ethics.
3. Orseau, L., & Armstrong, S. (2016). Safely Interruptible Agents. Proceedings of the 32nd Conference on Uncertainty in Artificial Intelligence, 557–566.
4. Pan, A., Bhatia, K., & Steinhardt, J. (2022). The Effects of Reward Misspecification: Mapping and Mitigating Misaligned Models. International Conference on Learning Representations.
5. Denison, C., MacDiarmid, M., Barez, F., Duvenaud, D., Kravec, S., Marks, S., Schiefer, N., Soklaski, R., Tamkin, A., Kaplan, J., Shlegeris, B., Bowman, S. R., & Perez, E. (2024). Sycophancy to Subterfuge: Investigating Reward-Tampering in Large Language Models. arXiv:2406.10162.
6. Hudson, R. (2026). Corrigibility Transformation: Constructing Goals That Accept Updates. arXiv:2510.15395.
7. Zhuang, S., & Hadfield-Menell, D. (2021). Consequences of Misaligned AI. Advances in Neural Information Processing Systems.
8. Carlsmith, J. (2022). Is Power-Seeking AI an Existential Risk? arXiv:2206.13353.
9. Christiano, P. F., Leike, J., Brown, T. B., Martic, M., Legg, S., & Amodei, D. (2017). Deep Reinforcement Learning from Human Preferences. Advances in Neural Information Processing Systems 30, 4299–4307.
10. Hadfield-Menell, D., Milli, S., Abbeel, P., Russell, S., & Dragan, A. (2017). Inverse Reward Design. Advances in Neural Information Processing Systems 30, 6765–6774.